

PARALLEL TOPOLOGICAL SORTING

DESIGN OF HIGH PERFORMANCE COMPUTING, FALL 2015


Kevin Wallimann Johannes Baum Matthias Untergassmair

ETH Zürich

December 14, 2015

OVERVIEW

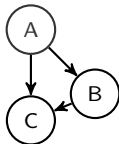
- DAG defines partial order
- Topological sorting defines one total order on a DAG
- Parallel algorithm: finds one topological sorting of a given DAG



DIFFERENCE TO BFS

- BFS visits every node
- Topological sorting algorithm needs to visit every edge

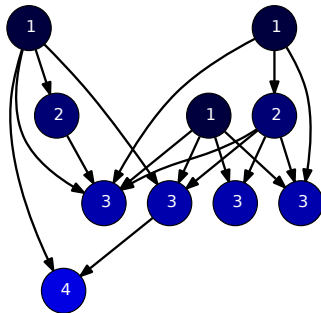
Example:



Consider order A,C,B \rightarrow valid in BFS, invalid in topological sorting

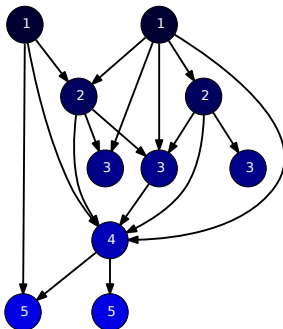
RANDOM GRAPH

- Parameter:
Average node
degree



SOFTWARE DEPENDENCY GRAPH

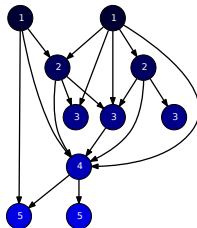
| Nodes | Edges | Degree (Median) |
|---------|---------|--------------------|
| 100'000 | 266'680 | 2 |
| 10'000 | 27'416 | 2 |



Musco, V. et al. (2014) "A Generative Model of Software Dependency Graphs to Better Understand Software Evolution."

PARALLEL ALGORITHM (SHARED MEMORY)

- ① As a preparing step, initialize a counter for every child with the number of its parents.
- ② Distribute parent nodes over threads and process them in parallel.
- ③ For every parent node, get a list of all child nodes and append the parent itself to solution (\rightarrow **lock**).
- ④ For every child of the parent, decrement its counter (\rightarrow **lock**). Once the counter is zero, we can move on
- ⑤ When all parents are processed (\rightarrow **barrier**), distribute new nodes and repeat



LOCAL LISTS APPROACH

- Idea:** perfect load balancing by redistributing nodes at every step
- Parent nodes stored in a global list.
 - Distribution of parent nodes: scatter the list among the threads. Each thread has now its local list.
 - Add new nodes to the end of local list.
 - When all parents were processed, gather all local lists into the global list.
 - Repeat until there are no parents left in the global list.

BOOLEAN ARRAY APPROACH

Idea: minimizing memory access by using lookup table

- Array of length N . 1 if node i is a parent node, 0 otherwise.
- Distribution of parent nodes: Parallel for-loop through the array.
- Mark new parents by setting a 1 in a second array.
- When all parents were processed, swap arrays.
- Loop through the array until there are no new parents.

OPTIMISTIC COUNTER CHECK

- Decrement shared counter
- Return true if counter is zero

```
inline bool counterCheck() {  
    bool lastone;  
    #pragma omp critical  
    {  
        --parcount_  
        lastone = (parcount_ == 0);  
    }  
    return lastone;  
}
```

OPTIMISTIC COUNTER CHECK

- Decrement shared counter
- Return true if counter is zero

```
inline bool counterCheck() {  
    #pragma omp atomic  
    --parcount_  
    return (parcount_ == 0);  
}
```

- Multiple threads could return true, although only one thread should do so.

OPTIMISTIC COUNTER CHECK

```
inline bool counterCheck() {  
    #pragma omp atomic  
    --parcount_  
    return (parcount_ == 0);  
}
```

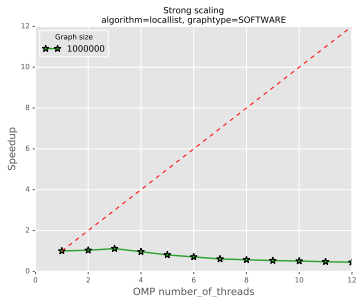
- List based approach: Child is inserted multiple times to solution \Rightarrow Wrong.
- Array approach: Multiple threads write 1 to the array \Rightarrow Ok, doesn't matter.

EULER

- Intel Xeon E5 on Euler cluster
- 2 processors per node
- 12 cores
- 30 MB shared last-level cache
- Software graph with around 1M nodes should fit into cache

STRONG SCALING SOFTWARE GRAPH

Local lists



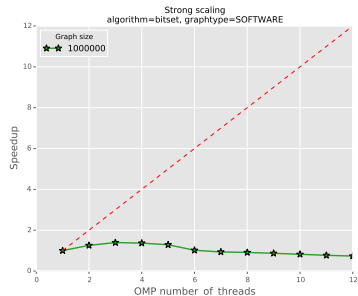
Absolute runtimes on 1 core

serial
0.45 s

bool array
0.58 s

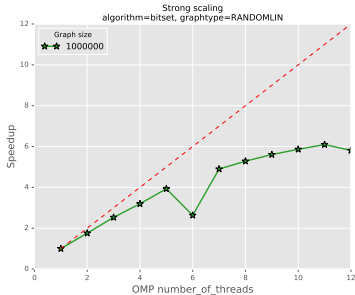
local list
0.48 s

Bool array



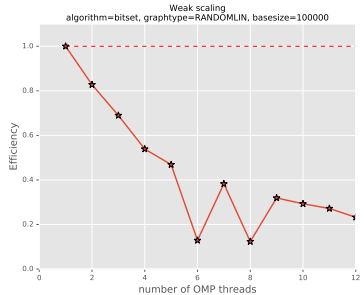
BOOL ARRAY & RANDOM GRAPH (NODE DEGREE 30)

Strong scaling



Absolute runtimes on 1 core

Weak scaling




serial
2.76 s


bool array
3.51 s

ABSOLUTE RUNTIME COMPOSITION

Local lists



Bool array



REMAINING ISSUES

- Pinpoint the reasons for bad scaling
- Work stealing could help to improve local list approach